

of A_0 and \mathcal{T} , and $d_0 \in A_0^\mathcal{I}$. For $d \in \Delta^\mathcal{I}$ and $i \leq \text{rd}(A_0)$, set

$$\text{tp}_i(d) = \{A \in \text{Def}_i(\mathcal{T}) \mid d \in A^\mathcal{I}\}.$$

We use \mathcal{I} to guide the nondeterministic decisions of the algorithm. To do this, it is convenient to pass an element $d \in \Delta^\mathcal{I}$ as a virtual fourth argument to the procedure `recurse` such that $d \in A^\mathcal{I}$ for all A in the first argument τ .

Initially, we guide the algorithm to guess $\text{tp}_{\text{rd}(A_0)}(d_0)$ as the set τ in \mathcal{ALC} -worlds. Now let `recurse` be called with arguments $(\tau, i, \mathcal{T}, d)$, and assume that the **for all** loop is processing $A \in \tau$ with $A \equiv \exists r.B \in \mathcal{T}$. Then $d \in A^\mathcal{I}$ and thus there is a $d' \in B^\mathcal{I}$ with $(d, d') \in r^\mathcal{I}$. We guide the algorithm to guess $\text{tp}_{i-1}(d')$ as the set τ . It remains to show that, when guided in this way, the algorithm returns **true**. This boils down to showing that all the guessed sets τ are types, which is straightforward using the semantics. \square

We have thus proved the following result.

Theorem 5.5. *In \mathcal{ALC} , concept satisfiability and subsumption with respect to acyclic TBoxes are in PSPACE.*

Lower Bound

We now prove that the PSPACE upper bound from Theorem 5.5 is optimal by showing that concept satisfiability in \mathcal{ALC} is PSPACE-hard, even without TBoxes. This implies that concept satisfiability is PSPACE-complete, both without TBoxes and with acyclic TBoxes.

The most common way to prove hardness for a complexity class \mathcal{C} is to find an appropriate problem P that is already known to be hard for \mathcal{C} and then to exhibit a polynomial time reduction from P to the problem at hand. In our case, the problem P is related to a game played on formulas of propositional logic and known to be PSPACE-complete [SC79].

A *finite Boolean game* (FBG) is a triple $(\varphi, \Gamma_1, \Gamma_2)$ with φ a formula of propositional logic and $\Gamma_1 \uplus \Gamma_2$ a partition of the variables used in φ into two sets of identical cardinality. The game is played by two players. Intuitively, Player 1 controls the variables in Γ_1 and Player 2 controls the variables in Γ_2 . The game proceeds in $n = |\Gamma_1 \uplus \Gamma_2|$ rounds, with the players alternating. We assume that the variables in Γ_1 and Γ_2 are ordered. Player 1 moves first by choosing a truth value for the first variable from Γ_1 . In the next round, Player 2 chooses a truth value for the first variable from Γ_2 . Next, it is again Player 1's turn, who assigns

a truth value to the second variable in Γ_1 , and so on. After n rounds, Player 1 wins the game if the resulting truth assignment satisfies the formula φ ; otherwise, Player 2 wins. The decision problem associated with this game is as follows: given a game $(\varphi, \Gamma_1, \Gamma_2)$, decide whether Player 1 has a *winning strategy*, i.e., whether he can force a win no matter what Player 2 does.

Before reducing FBGs to concept satisfiability in \mathcal{ALC} , we give a formal definition of winning strategies. Fix a game $G = (\varphi, \Gamma_1, \Gamma_2)$, and let $n = |\Gamma_1 \uplus \Gamma_2|$. We assume that $\Gamma_1 = \{p_1, p_3, \dots, p_{n-1}\}$ and $\Gamma_2 = \{p_2, p_4, \dots, p_n\}$. A *configuration* of G is a word $t \in \{0, 1\}^i$, for some $i \leq n$. Intuitively, the k th symbol in this word assigns a truth value to the variable p_k . Thus, if the current configuration is t , then a truth value for $p_{|t|+1}$ is selected in the next round. This is done by Player 1 if $|t|$ is even and by Player 2 if $|t|$ is odd. The *initial configuration* is the empty word ε . A *winning strategy* for Player 1 in G is a finite node-labelled tree (V, E, ℓ) of depth n , where ℓ assigns to each node $v \in V$ a configuration $\ell(v)$. We say that a node $v \in V$ is of *depth* i if v is reachable from the root by travelling along i successor edges. Winning strategies are required to satisfy the following conditions:


- the root is labelled with the initial configuration;
- if v is a node of depth $i < n$ with i even and $\ell(v) = t$, then v has one successor v' with $\ell(v') \in \{t0, t1\}$;
- if v is a node of depth $i < n$ with i odd and $\ell(v) = t$, then v has two successors v' and v'' with $\ell(v') = t0$ and $\ell(v'') = t1$;
- if v is a node of depth n and $\ell(v) = t$, then t satisfies φ .

Consider the game $G = (\varphi, \{p_1, p_3\}, \{p_2, p_4\})$, with

$$\varphi = (\neg p_1 \rightarrow p_2) \wedge ((p_1 \wedge p_2) \rightarrow (p_3 \vee p_4)) \wedge (\neg p_2 \rightarrow (p_4 \rightarrow \neg p_3)).$$

Figure 5.2 shows a winning strategy for Player 1 in G . Intuitively, a winning strategy tells Player 1 how to play the game, no matter what Player 2 does. For example, if the current game configuration is 10, then Player 1 can look into the strategy tree for the (unique!) node $v \in V$ with $\ell(v) = 10$ and at its (unique!) successor v' . It is labelled 100, which advises Player 1 to set the truth value of p_3 to 0.

To reduce the existence of winning strategies in FBGs to the satisfiability of \mathcal{ALC} concepts, we transform a game $G = (\varphi, \Gamma_1, \Gamma_2)$ into an \mathcal{ALC} concept C_G such that Player 1 has a winning strategy in G if and only if C_G is satisfiable. The idea is to craft C_G such that every model

Fig. 5.2. A winning strategy for Player 1 in G .

of C_G describes a winning strategy for Player 1 in G and, vice versa, every such winning strategy gives rise to a model of C_G . The concept C_G uses a single role name r to represent the successor relation of the strategy tree. To describe the values of propositional variables, we use concept names P_1, \dots, P_n . Throughout this chapter, we use $C \rightarrow D$ as an abbreviation for $\neg C \sqcup D$, for better readability. Now, C_G is a conjunction whose conjuncts we define step by step, along with intuitive explanations of their meaning:

- For each node of odd depth i (i.e., Player 2 is to move), there are two successors, one for each possible truth value of p_{i+1} :

$$C_1 = \bigcap_{i \in \{1, 3, \dots, n-1\}} \forall r^i. (\exists r. \neg P_{i+1} \sqcap \exists r. P_{i+1}),$$

where $\forall r^i.C$ denotes the i -fold nesting $\forall r. \dots \forall r.C$.

- For each node of even depth i (i.e., Player 1 is to move), there is one successor:

$$C_2 = \bigcap_{i \in \{0, 2, \dots, n-2\}} \forall r^i. \exists r. \top.$$

Note that, since P_{i+1} must be either true or false at the generated successor, a truth value for p_{i+1} is chosen “automatically”.

- Once a truth value is chosen, it remains fixed:

$$C_3 = \bigcap_{1 \leq i \leq j < n} \forall r^j. ((P_i \rightarrow \forall r. P_i) \sqcap (\neg P_i \rightarrow \forall r. \neg P_i)).$$

- At the leaves, the formula φ is true:

$$C_4 = \forall r^n. \varphi^*,$$

where φ^* denotes the result of converting φ into an \mathcal{ALC} concept by replacing each p_i with P_i , \sqcap with \wedge , and \sqcup with \vee .

Now, we define $C_G = C_1 \sqcap \dots \sqcap C_4$. It is easily verified that the length of C_G is quadratic in n , and that C_G can be constructed in time polynomial in n . The next lemma states that the reduction is correct.

Lemma 5.6. *Player 1 has a winning strategy in G if and only if C_G is satisfiable.*

Proof. (only if) Assume that Player 1 has a winning strategy (V, E, ℓ) with root $v_0 \in V$. We define an interpretation \mathcal{I} by setting

- $\Delta^{\mathcal{I}} = V$,
- $r^{\mathcal{I}} = E$,
- $P_i^{\mathcal{I}} = \{v \in V \mid \ell(v) \text{ sets } p_i \text{ to } 1\}$ for $1 \leq i \leq n$.

We leave it as an exercise to verify that $v_0 \in C_G^{\mathcal{I}}$.

(if) Let \mathcal{I} be a model of C_G , and let $d_0 \in C_G^{\mathcal{I}}$. We define a winning strategy (V, E, ℓ) with $V \subseteq \mathbb{N} \times \Delta^{\mathcal{I}}$. The construction will be such that

- (*) if $(i, d) \in V$, then d is reachable from d_0 in \mathcal{I} by travelling i steps along r .

Start by setting $V = \{(0, d_0)\}$, $E = \emptyset$ and $\ell(0, d_0)$ to the initial configuration. We proceed in rounds $1, \dots, n$. In each odd round i , iterate over all nodes $(i-1, d) \in V$ and do the following:

- select a $d' \in \Delta^{\mathcal{I}}$ such that $(d, d') \in r^{\mathcal{I}}$ (which exists since $d_0 \in C_2^{\mathcal{I}}$ and (*) is satisfied by induction);
- add (i, d') to V , $((i-1, d), (i, d'))$ to E , and set $\ell(i, d') = tj$, where $t = \ell(i-1, d)$ and j is 1 if $d' \in P_i^{\mathcal{I}}$ and 0 otherwise.

In each even round i , iterate over all nodes $(i-1, d) \in V$ and do the following:

- select $d', d'' \in \Delta^{\mathcal{I}}$ such that $d' \in P_i^{\mathcal{I}}$, $d'' \notin P_i^{\mathcal{I}}$ and $\{(d, d'), (d, d'')\} \subseteq r^{\mathcal{I}}$ (which exist since $d_0 \in C_1^{\mathcal{I}}$ and (*) is satisfied by induction);
- add (i, d') and (i, d'') to V , $((i-1, d), (i, d'))$ and $((i-1, d), (i, d''))$ to E , and set $\ell(i, d') = t1$ and $\ell(i, d'') = t0$, where $t = \ell(i-1, d)$.

Since $d_0 \in C_3^{\mathcal{I}}$ and $d_0 \in C_4^{\mathcal{I}}$, it is easy to prove that the resulting tree is a winning strategy for Player 1. \square

We have thus established PSPACE-hardness of concept satisfiability in \mathcal{ALC} . Together with Theorem 5.5, we obtain the following result.

Theorem 5.7. *In \mathcal{ALC} , concept satisfiability and subsumption without TBoxes and with acyclic TBoxes are PSPACE-hard, thus PSPACE-complete.*

As in the case of \mathcal{ALC} , it is rather often the case that satisfiability without TBoxes and with acyclic TBoxes have the same complexity. However, there are also notable exceptions. One example is \mathcal{ALC} extended with concrete domains. For some natural concrete domains, satisfiability without TBoxes is PSPACE-complete, and with respect to acyclic TBoxes it is NEXPTIME-complete.

5.1.2 General TBoxes

The aim of this section is to show that, in \mathcal{ALC} , the transition from acyclic TBoxes to general TBoxes increases the computational complexity from PSPACE to EXPTIME.

Upper Bound

We prove an EXPTIME upper bound for satisfiability with respect to general \mathcal{ALC} concepts using a so-called *type elimination* algorithm. The central notion of such an algorithm is that of a type, which is defined similarly to the types introduced in Section 5.1.1.

Let \mathcal{T} be a general TBox. It can be seen that \mathcal{T} is equivalent to the TBox

$$\top \sqsubseteq \bigcap_{C \sqsubseteq D \in \mathcal{T}} \neg C \sqcup D :$$

see Point (v) of Lemma 2.16 for a similar observation. We can thus assume without loss of generality that general TBoxes \mathcal{T} have the form $\{\top \sqsubseteq C_{\mathcal{T}}\}$. Moreover, we can assume that $C_{\mathcal{T}}$ is in negation normal form (NNF); compare Chapter 4. As in Section 3.4, we use $\text{sub}(C)$ to denote the set of subconcepts of the concept C . If \mathcal{T} is a TBox, we set $\text{sub}(\mathcal{T}) = \text{sub}(C_{\mathcal{T}})$.

Definition 5.8. Let \mathcal{T} be a general TBox. A *type* for \mathcal{T} is a set $\tau \subseteq \text{sub}(\mathcal{T})$ such that the following conditions are satisfied: